

Final Report

November 25, 2022

1 2D Design Template

2 Overview

The purpose of this project is for you to apply what you have learnt in this course. This includes working with data and visualizing it, create model of linear regression or logistic regression, as well as using metrics to measure the accuracy of your model.

Please find the project handout description in the following link: - [DDW-MU-Humanities Handout](#) - [DDW-MU-SocialStudies Handout](#)

2.1 Deliverables

You need to submit this Jupyter notebook together with the dataset into Vocareum. Use the template in this notebook to work on this project.

2.2 Students Submission

Student's Name: - Bundhoo Simriti - Elvern Neylmav Tanny - Koh Chee Kiat - Haritha Shraeya Rajasekar - Mahima Sharma - Zhang Jianyu

```
In [1]: # Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

2.3 Overview About the Problem

Describe here the problem you are trying to solve.

2.3.1 Problem Statement

We aim to predict the future of food safety and security in countries of the lower income bracket like Cambodia and Myanmar through the prediction of amount of undernourishment as it encapsulates percentage of people whose food requirements are not satisfied.

- Southeast Asia is a diverse, fast-growing region, making remarkable progress in terms of improving food security, going from 31% undernourishment in the 1990s to below 10% by 2014-2016 (<https://www.fao.org/3/bt099e/bt099e.pdf>). Although undernourishment has been decreasing, food security is still a concern when accounting for the fast-growing population in Southeast Asia, projected to grow from 640 million to more than 710 million by 2030 (<https://www.enterprisesg.gov.sg/overseas-markets/asia-pacific/asean/overview>). Therefore, our group focused on modelling the food security of Southeast Asia.

2.4 Dataset

Describe here your data set. Put the link to the sources of your dataset. Describe your data and what are the columns.

Put some Python codes here to describe and visualize your data.

Description of Dataset: The GDP per capita dataset documents gross domestic product per person in USD of a country since 1960.

Inflation, consumer prices (annual %) documents the increase or decrease in inflation rate in comparison to the previous year since 1960.

Gross per capita Production Index Number (2014-2016 = 100) (Food production index) documents relative index of food production in the country. With the 3 year period 2014-2016 given 100 points as a reference since 2001.

Percentage of Undernourishment Prevalence documents percentage of people whose food requirements aren't satisfied since 2000.

Our model uses 'GDP per Capita (USD)', 'Annual Inflation Rate (%)', 'Food production index (2014-2016 = 100)' as features and 'Percentage of Undernourishment Prevalence' as target spanning from 2001 to 2020

Sources: GDP per capita dataset (USD) - <https://data.worldbank.org/indicator/NY.GDP.PCAP.CD>

Inflation, consumer prices (annual %) - <https://data.worldbank.org/indicator/FP.CPI.TOTL.ZG>

Gross per capita Production Index Number (2014-2016 = 100) (Food production index) - <https://www.fao.org/faostat/en/#data/QI>

Percentage of Undernourishment Prevalence - <https://www.fao.org/faostat/en/#data/FS>

Dataset iterations: We first focused on all countries in Southeast Asia and then worked our way to decreasing the number of countries as the data range was too diverse and hence meaningful conclusions could not be achieved. We focussed our scope to emphasize on countries with lower-income and having agricultural significance, specifically Myanmar and Cambodia.

We decided to narrow our scope to lower-middle income countries categorised by the World Bank (<https://datatopics.worldbank.org/world-development-indicators/the-world-by-income-and-region.html>) and focus on countries whose economies are more dependent on agriculture as such countries are more at risk of facing food security issues. This left us with Cambodia and Myanmar, where agriculture makes up more than 20% of their GDP (22.7% and 20.9% respectively, 2020) <https://data.worldbank.org/indicator/NV.AGR.TOTL.ZS?end=2020&locations=KH-MM&start=2020&view=bar>

We explored multiple categories of features and performed linear regression on them one at a time and identified those features with which we were able to identify a relation: economic,

environmental, food production features. After testing them, we came to a conclusion to use the following economic measures 'GDP per Capita (USD)', 'Annual Inflation Rate (%)', 'Food production index (2014-2016 = 100)' as our features. We chose prevalence of undernourishment as a (%) as our target.

2.5 Cleaning Dataset

2.5.1 Percentage of Undernourishment Prevalence

In [2]: # Read the Percentage of Undernourishment Prevalence CSV file online

```
df_pup_url = "https://raw.githubusercontent.com/verneylmavt/2D_Project_Term-3/main/2D_I
df_pup = pd.read_csv(df_pup_url, encoding='latin-1')
df_pup = df_pup.set_index("Area")
display(df_pup)
```

Area	iz&Domain Code	Domain	\
Afghanistan	FS Suite of Food Security Indicators		
Afghanistan	FS Suite of Food Security Indicators		
Afghanistan	FS Suite of Food Security Indicators		
Afghanistan	FS Suite of Food Security Indicators		
Afghanistan	FS Suite of Food Security Indicators		
...	
Zimbabwe	FS Suite of Food Security Indicators		
Zimbabwe	FS Suite of Food Security Indicators		
Zimbabwe	FS Suite of Food Security Indicators		
Zimbabwe	FS Suite of Food Security Indicators		
Zimbabwe	FS Suite of Food Security Indicators		

Area	Area Code (M49)	Element Code	Element	Item Code	\
Afghanistan	4	6121	Value	210041	
Afghanistan	4	6121	Value	210041	
Afghanistan	4	6121	Value	210041	
Afghanistan	4	6121	Value	210041	
Afghanistan	4	6121	Value	210041	
...	
Zimbabwe	716	6121	Value	210041	
Zimbabwe	716	6121	Value	210041	
Zimbabwe	716	6121	Value	210041	
Zimbabwe	716	6121	Value	210041	
Zimbabwe	716	6121	Value	210041	

Area	Item	Year	Code	\
Afghanistan	Prevalence of undernourishment (percent)	(3-ye...	20002002	
Afghanistan	Prevalence of undernourishment (percent)	(3-ye...	20012003	
Afghanistan	Prevalence of undernourishment (percent)	(3-ye...	20022004	
Afghanistan	Prevalence of undernourishment (percent)	(3-ye...	20032005	

Afghanistan	Prevalence of undernourishment (percent) (3-ye...	20042006
...
Zimbabwe	Prevalence of undernourishment (percent) (3-ye...	20152017
Zimbabwe	Prevalence of undernourishment (percent) (3-ye...	20162018
Zimbabwe	Prevalence of undernourishment (percent) (3-ye...	20172019
Zimbabwe	Prevalence of undernourishment (percent) (3-ye...	20182020
Zimbabwe	Prevalence of undernourishment (percent) (3-ye...	20192021

	Year	Unit	Value	Flag	Flag	Description	Note
Area							
Afghanistan	2000-2002	%	47.8	E		Estimated value	NaN
Afghanistan	2001-2003	%	45.6	E		Estimated value	NaN
Afghanistan	2002-2004	%	40.6	E		Estimated value	NaN
Afghanistan	2003-2005	%	38	E		Estimated value	NaN
Afghanistan	2004-2006	%	36.1	E		Estimated value	NaN
...
Zimbabwe	2015-2017	%	NaN	0		Missing value	NaN
Zimbabwe	2016-2018	%	NaN	0		Missing value	NaN
Zimbabwe	2017-2019	%	NaN	0		Missing value	NaN
Zimbabwe	2018-2020	%	NaN	0		Missing value	NaN
Zimbabwe	2019-2021	%	NaN	0		Missing value	NaN

[4102 rows x 14 columns]

```
In [3]: # Extracting data for Myanmar and Cambodia
df_pup_cambodia_myanmar = (df_pup.loc[["Cambodia", "Myanmar"], :]).copy()
cambodia_pup = (df_pup_cambodia_myanmar.loc["Cambodia", "Value"]).tolist()
myanmar_pup = (df_pup_cambodia_myanmar.loc["Myanmar", "Value"]).tolist()
```

2.5.2 GDP per Capita (USD)

```
In [4]: # Read the GDP per capita (USD) CSV file online
df_gdp_url = "https://raw.githubusercontent.com/verneylmavt/2D_Project_Term-3/main/2D_1
df_gdp = pd.read_csv(df_gdp_url)
df_gdp = df_gdp.set_index("Country Name")
display(df_gdp)
```

Country Name	Country Code	Indicator Name \
Aruba	ABW	GDP per capita (current US\$)
Africa Eastern and Southern	AFE	GDP per capita (current US\$)
Afghanistan	AFG	GDP per capita (current US\$)
Africa Western and Central	AFW	GDP per capita (current US\$)
Angola	AGO	GDP per capita (current US\$)
...
Kosovo	KKX	GDP per capita (current US\$)
Yemen, Rep.	YEM	GDP per capita (current US\$)

South Africa	ZAF	GDP per capita (current US\$)
Zambia	ZMB	GDP per capita (current US\$)
Zimbabwe	ZWE	GDP per capita (current US\$)

Country Name	Indicator Code	1960	1961 \
Aruba	NY.GDP.PCAP.CD	NaN	NaN
Africa Eastern and Southern	NY.GDP.PCAP.CD	162.726326	162.555968
Afghanistan	NY.GDP.PCAP.CD	59.773234	59.860900
Africa Western and Central	NY.GDP.PCAP.CD	107.930722	113.080062
Angola	NY.GDP.PCAP.CD	NaN	NaN
...
Kosovo	NY.GDP.PCAP.CD	NaN	NaN
Yemen, Rep.	NY.GDP.PCAP.CD	NaN	NaN
South Africa	NY.GDP.PCAP.CD	511.618737	526.461750
Zambia	NY.GDP.PCAP.CD	232.188565	220.042067
Zimbabwe	NY.GDP.PCAP.CD	278.813847	280.828663

Country Name	1962	1963	1964	1965 \
Aruba	NaN	NaN	NaN	NaN
Africa Eastern and Southern	172.271022	199.784916	180.228774	199.517227
Afghanistan	58.458009	78.706429	82.095307	101.108325
Africa Western and Central	118.829461	123.441090	131.852423	138.524029
Angola	NaN	NaN	NaN	NaN
...
Kosovo	NaN	NaN	NaN	NaN
Yemen, Rep.	NaN	NaN	NaN	NaN
South Africa	546.261935	589.160460	632.716104	674.186433
Zambia	212.578449	213.896759	242.384472	303.281741
Zimbabwe	276.688233	277.479715	281.558896	293.308788

Country Name	1966	...	2012	2013 \
Aruba	NaN	...	25496.843940	26442.426800
Africa Eastern and Southern	211.054388	...	1777.303950	1748.905594
Afghanistan	137.594298	...	638.845852	624.315454
Africa Western and Central	144.323882	...	1965.115750	2157.494584
Angola	NaN	...	4978.434435	5127.717243
...
Kosovo	NaN	...	3410.859780	3704.784221
Yemen, Rep.	NaN	...	1446.536472	1607.152173
South Africa	714.562010	...	8222.197279	7467.079185
Zambia	343.373670	...	1763.069442	1878.346811
Zimbabwe	277.234532	...	1304.968011	1429.998461

Country Name	2014	2015	2016 \
--------------	------	------	--------

Aruba	26895.057170	28399.050130	28453.715560
Africa Eastern and Southern	1736.242220	1556.316469	1446.533624
Afghanistan	614.223342	556.007221	512.012778
Africa Western and Central	2212.914095	1894.322115	1673.843681
Angola	5094.112329	3127.890598	1728.023754
...
Kosovo	3902.676013	3520.766449	3759.560246
Yemen, Rep.	1674.002572	1601.807163	1152.738019
South Africa	6988.808739	6259.839681	5756.965741
Zambia	1762.427817	1338.290927	1280.806543
Zimbabwe	1434.896277	1445.069702	1464.588957

	2017	2018	2019 \
Country Name			
Aruba	29348.418970	30253.714230	31135.884360
Africa Eastern and Southern	1629.404273	1541.031661	1511.309259
Afghanistan	516.679862	485.668419	494.179350
Africa Western and Central	1613.490478	1704.135698	1777.852822
Angola	2313.220584	2524.942483	2177.799015
...
Kosovo	4009.380987	4384.048892	4416.108358
Yemen, Rep.	964.340344	758.145949	750.554583
South Africa	6690.939847	7005.095413	6624.761865
Zambia	1535.196574	1516.368371	1305.001031
Zimbabwe	1235.189032	1254.642265	1316.740657

	2020	2021
Country Name		
Aruba	23384.298790	NaN
Africa Eastern and Southern	1360.878645	1557.722682
Afghanistan	516.747871	NaN
Africa Western and Central	1709.764129	1774.921218
Angola	1631.431691	2137.909393
...
Kosovo	4310.811183	4986.582469
Yemen, Rep.	631.681490	690.759273
South Africa	5655.867654	6994.211654
Zambia	985.132436	1120.630171
Zimbabwe	1214.509820	1737.173977

[266 rows x 65 columns]

```
In [5]: # Extracting data for Myanmar and Cambodia for required year range
df_gdp_cambodia_myanmar = (df_gdp.loc[["Cambodia", "Myanmar"], "2001":"2020"]).copy()
cambodia_gdp = (df_gdp_cambodia_myanmar.loc["Cambodia", "2001":"2020"]).tolist()
myanmar_gdp = (df_gdp_cambodia_myanmar.loc["Myanmar", "2001":"2020"]).tolist()
```

2.5.3 Annual Inflation Rate (%)

In [6]: # Read the Annual Inflation Rate (%) CSV file online

```
df_inflation_rate_url = "https://raw.githubusercontent.com/verneylmavt/2D_Project_Term
df_inflation_rate = pd.read_csv(df_inflation_rate_url)
df_inflation_rate = df_inflation_rate.set_index("Country Name")
display(df_inflation_rate)
```

Country Name	Country Code \
Aruba	ABW
Africa Eastern and Southern	AFE
Afghanistan	AFG
Africa Western and Central	AFW
Angola	AGO
...	...
Kosovo	XKX
Yemen, Rep.	YEM
South Africa	ZAF
Zambia	ZMB
Zimbabwe	ZWE

Country Name	Indicator Name \
Aruba	Inflation, consumer prices (annual %)
Africa Eastern and Southern	Inflation, consumer prices (annual %)
Afghanistan	Inflation, consumer prices (annual %)
Africa Western and Central	Inflation, consumer prices (annual %)
Angola	Inflation, consumer prices (annual %)
...	...
Kosovo	Inflation, consumer prices (annual %)
Yemen, Rep.	Inflation, consumer prices (annual %)
South Africa	Inflation, consumer prices (annual %)
Zambia	Inflation, consumer prices (annual %)
Zimbabwe	Inflation, consumer prices (annual %)

Country Name	Indicator Code	1960	1961	1962 \
Aruba	FP.CPI.TOTL.ZG	NaN	NaN	NaN
Africa Eastern and Southern	FP.CPI.TOTL.ZG	NaN	NaN	NaN
Afghanistan	FP.CPI.TOTL.ZG	NaN	NaN	NaN
Africa Western and Central	FP.CPI.TOTL.ZG	NaN	NaN	NaN
Angola	FP.CPI.TOTL.ZG	NaN	NaN	NaN
...
Kosovo	FP.CPI.TOTL.ZG	NaN	NaN	NaN
Yemen, Rep.	FP.CPI.TOTL.ZG	NaN	NaN	NaN
South Africa	FP.CPI.TOTL.ZG	1.288859	2.102374	1.246285
Zambia	FP.CPI.TOTL.ZG	NaN	NaN	NaN
Zimbabwe	FP.CPI.TOTL.ZG	NaN	NaN	NaN

	1963	1964	1965	1966	...	\
Country Name					...	
Aruba	NaN	NaN	NaN	NaN	...	
Africa Eastern and Southern	NaN	NaN	NaN	NaN	...	
Afghanistan	NaN	NaN	NaN	NaN	...	
Africa Western and Central	NaN	NaN	NaN	NaN	...	
Angola	NaN	NaN	NaN	NaN	...	
...	
Kosovo	NaN	NaN	NaN	NaN	...	
Yemen, Rep.	NaN	NaN	NaN	NaN	...	
South Africa	1.33797	2.534972841	4.069029	3.489234	...	
Zambia	NaN	NaN	NaN	NaN	...	
Zimbabwe	NaN	NaN	NaN	NaN	...	

	2012	2013	2014	2015	\
Country Name					
Aruba	0.627472	-2.372065	0.421441	0.474764	
Africa Eastern and Southern	9.158707	5.750981	5.370290	5.250171	
Afghanistan	6.441213	7.385772	4.673996	-0.661709	
Africa Western and Central	4.578375	2.439201	1.758052	2.130268	
Angola	10.277905	8.777814	7.280387	9.150372	
...	
Kosovo	2.476738	1.767324	0.428958	-0.536929	
Yemen, Rep.	9.885387	10.968442	8.104726	NaN	
South Africa	5.724658	5.784469	6.129838	4.540642	
Zambia	6.575900	6.977676	7.806876	10.110593	
Zimbabwe	3.725327	1.634950	-0.197785	-2.430968	

	2016	2017	2018	2019	\
Country Name					
Aruba	-0.931196	-1.028282	3.626041	4.257462	
Africa Eastern and Southern	6.571396	6.399343	4.720811	4.120246	
Afghanistan	4.383892	4.975952	0.626149	2.302373	
Africa Western and Central	1.494564	1.764635	1.784050	1.758565	
Angola	30.695313	29.843587	19.628608	17.081215	
...	
Kosovo	0.273169	1.488234	1.053798	2.675992	
Yemen, Rep.	NaN	NaN	NaN	NaN	
South Africa	6.571396	5.184247	4.517165	4.120246	
Zambia	17.869730	6.577312	7.494572	9.150316	
Zimbabwe	-1.543670	0.893962	10.618866	255.304991	

	2020	2021
Country Name		
Aruba	NaN	NaN
Africa Eastern and Southern	5.404815	7.240978
Afghanistan	NaN	NaN

Africa Western and Central	2.492522	3.925603
Angola	NaN	NaN
...
Kosovo	0.198228	3.353691
Yemen, Rep.	NaN	NaN
South Africa	3.210036	4.611672
Zambia	15.732585	22.021234
Zimbabwe	557.201817	98.546105

[266 rows x 65 columns]

```
In [7]: # Extracting data for Myanmar and Cambodia for required year range
df_inflation_rate_cambodia_myanmar = (df_inflation_rate.loc[["Cambodia", "Myanmar"], "2001": "2016"])
cambodia_inflation_rate = (df_inflation_rate_cambodia_myanmar.loc["Cambodia", "2001": "2016"])
myanmar_inflation_rate = (df_inflation_rate_cambodia_myanmar.loc["Myanmar", "2001": "2016"])
```

2.5.4 Gross per capita Production Index Number (2014-2016 = 100)

```
In [8]: # Read the Production Index Number CSV file online
df_pin1416_url = "https://raw.githubusercontent.com/verneylmavt/2D_Project_Term-3/main/Production_Index_Number.csv"
df_pin1416 = pd.read_csv(df_pin1416_url, encoding='latin-1')
df_pin1416 = df_pin1416.set_index("Area")
display(df_pin1416)
```

Area	izfDomain Code	Domain	Area Code (M49)	Element Code \
Afghanistan	QI	Production Indices	4	434
Afghanistan	QI	Production Indices	4	434
Afghanistan	QI	Production Indices	4	434
Afghanistan	QI	Production Indices	4	434
Afghanistan	QI	Production Indices	4	434
...
Zimbabwe	QI	Production Indices	716	434
Zimbabwe	QI	Production Indices	716	434
Zimbabwe	QI	Production Indices	716	434
Zimbabwe	QI	Production Indices	716	434
Zimbabwe	QI	Production Indices	716	434

Area	Element \
Afghanistan	Gross per capita Production Index Number (2014-2016 = 100)
Afghanistan	Gross per capita Production Index Number (2014-2016 = 100)
Afghanistan	Gross per capita Production Index Number (2014-2016 = 100)
Afghanistan	Gross per capita Production Index Number (2014-2016 = 100)
Afghanistan	Gross per capita Production Index Number (2014-2016 = 100)
...	...
Zimbabwe	Gross per capita Production Index Number (2014-2016 = 100)

Zimbabwe	Gross per capita Production Index Number (2014...
Zimbabwe	Gross per capita Production Index Number (2014...
Zimbabwe	Gross per capita Production Index Number (2014...
Zimbabwe	Gross per capita Production Index Number (2014...

	Item Code (CPC)	Item	Year	Code	Year	Unit	Value	Flag	\
Area									
Afghanistan	F2051	Agriculture	1961	1961	index	161.93	E		
Afghanistan	F2051	Agriculture	1962	1962	index	161.58	E		
Afghanistan	F2051	Agriculture	1963	1963	index	159.56	E		
Afghanistan	F2051	Agriculture	1964	1964	index	166.81	E		
Afghanistan	F2051	Agriculture	1965	1965	index	170.37	E		
...
Zimbabwe	F2051	Agriculture	2016	2016	index	95.59	E		
Zimbabwe	F2051	Agriculture	2017	2017	index	100.70	E		
Zimbabwe	F2051	Agriculture	2018	2018	index	115.99	E		
Zimbabwe	F2051	Agriculture	2019	2019	index	91.59	E		
Zimbabwe	F2051	Agriculture	2020	2020	index	105.22	E		

	Flag Description
Area	
Afghanistan	Estimated value
Afghanistan	Estimated value
Afghanistan	Estimated value
Afghanistan	Estimated value
Afghanistan	Estimated value
...	...
Zimbabwe	Estimated value
Zimbabwe	Estimated value
Zimbabwe	Estimated value
Zimbabwe	Estimated value
Zimbabwe	Estimated value

[10920 rows x 13 columns]

```
In [9]: # Extracting data for Myanmar and Cambodia for required year range
df_pin1416_cambodia_myanmar = (df_pin1416.loc[["Cambodia", "Myanmar"], :]).copy()
df_pin1416_cambodia_myanmar["Year Code"] = df_pin1416_cambodia_myanmar["Year Code"].astype(int)
df_pin1416_cambodia_myanmar = df_pin1416_cambodia_myanmar.loc[
    (df_pin1416_cambodia_myanmar["Year Code"] >= 2001) & (df_pin1416_cambodia_myanmar["Year Code"] <= 2020)]

cambodia_pin1416 = (df_pin1416_cambodia_myanmar.loc["Cambodia", "Value"]).tolist()
myanmar_pin1416 = (df_pin1416_cambodia_myanmar.loc["Myanmar", "Value"]).tolist()
```

2.6 Combining Relevant Data Extracted into DataFrame

Note: With datasets using three year average, we considered the central year. Eg: 2000-2002 => 2001

```
In [10]: #Instantiating a new DataFrame called df, and adding a new column correspondingly
df = pd.DataFrame()
df["Country"] = pd.concat([pd.DataFrame(np.full((20,), "Cambodia")), pd.DataFrame(np.full((20,), "Myanmar"))], axis=0)
df["Year"] = df_pin1416_cambodia_myanmar.loc[:, "Year Code"].tolist()
df["Percentage of Undernourishment Prevalence (3-Year Average)"] = cambodia_pup + myanmar_pup
df["Binary Categorical"] = pd.concat([pd.DataFrame(np.full((20,), 0)), pd.DataFrame(np.full((20,), 1))], axis=0)
df["GDP per Capita (USD)"] = cambodia_gdp + myanmar_gdp
df["Annual Inflation Rate (%)"] = cambodia_inflation_rate + myanmar_inflation_rate
df["Food Production Index (2014-2016 = 100)"] = cambodia_pin1416 + myanmar_pin1416
```

```
In [11]: # 7 Columns: 4 Features & 1 Targets with 2 additional columns as a Description (Country & Year)
# 40 Rows, 2001-2020 Cambodia (20) & 2001-2020 Myanmar (20)
display(df)
print(df.shape)
```

	Country	Year	Percentage of Undernourishment Prevalence (3-Year Average) \
0	Cambodia	2001	23.6
1	Cambodia	2002	21.2
2	Cambodia	2003	19.4
3	Cambodia	2004	18.5
4	Cambodia	2005	17
5	Cambodia	2006	15.6
6	Cambodia	2007	14.8
7	Cambodia	2008	14.5
8	Cambodia	2009	13
9	Cambodia	2010	11.2
10	Cambodia	2011	9.7
11	Cambodia	2012	9.5
12	Cambodia	2013	9.4
13	Cambodia	2014	9.2
14	Cambodia	2015	8.9
15	Cambodia	2016	8.5
16	Cambodia	2017	7.7
17	Cambodia	2018	6.6
18	Cambodia	2019	6
19	Cambodia	2020	6.3
20	Myanmar	2001	37.6
21	Myanmar	2002	34.8
22	Myanmar	2003	32.4
23	Myanmar	2004	30.2
24	Myanmar	2005	27.8
25	Myanmar	2006	24.5
26	Myanmar	2007	20.5
27	Myanmar	2008	17.1

28	Myanmar	2009	12.9
29	Myanmar	2010	10.2
30	Myanmar	2011	7.8
31	Myanmar	2012	7.1
32	Myanmar	2013	6.1
33	Myanmar	2014	5.1
34	Myanmar	2015	4.2
35	Myanmar	2016	3.5
36	Myanmar	2017	3
37	Myanmar	2018	2.6
38	Myanmar	2019	<2.5
39	Myanmar	2020	3.1

	Binary	Categorical	GDP per Capita (USD)	Annual Inflation Rate (%)	\
0		0	321.150224	-0.600648	
1		0	338.987477	0.211467	
2		0	362.335482	0.941746	
3		0	408.513639	4.319337	
4		0	474.111192	6.615259	
5		0	539.750329	5.810686	
6		0	631.525258	8.708828	
7		0	745.609127	24.096852	
8		0	738.054731	-1.241718	
9		0	785.502667	3.996395	
10		0	882.275614	5.478447	
11		0	950.880346	2.934316	
12		0	1013.420536	2.941625	
13		0	1093.495976	3.855689	
14		0	1162.904995	1.223932	
15		0	1269.591499	3.019140	
16		0	1385.260066	2.912636	
17		0	1512.126989	2.459085	
18		0	1643.121389	1.942575	
19		0	1547.511388	2.940295	
20		1	131.715298	21.101305	
21		1	128.099702	57.074511	
22		1	161.055524	36.589718	
23		1	193.368766	4.534214	
24		1	216.311501	9.368618	
25		1	240.624014	19.996487	
26		1	314.202294	35.024597	
27		1	460.908889	26.799537	
28		1	586.168180	1.472343	
29		1	746.945360	7.718382	
30		1	1061.344429	5.021460	
31		1	1134.302224	1.467583	
32		1	1168.165453	5.643039	
33		1	1210.097654	4.953299	

34	1	1196.743333	9.454172
35	1	1136.610627	6.928825
36	1	1151.114464	4.572537
37	1	1250.173685	6.872329
38	1	1271.111536	8.825067
39	1	1450.662673	7.092387

Food Production Index (2014-2016 = 100)

0	52.79
1	49.46
2	57.65
3	53.83
4	68.71
5	73.50
6	75.90
7	81.68
8	84.33
9	89.31
10	100.67
11	102.56
12	102.11
13	99.42
14	98.42
15	102.17
16	105.05
17	106.81
18	106.21
19	103.61
20	59.21
21	60.64
22	65.22
23	69.95
24	77.51
25	86.08
26	90.42
27	96.56
28	100.13
29	102.32
30	97.60
31	95.65
32	99.23
33	98.77
34	101.00
35	100.22
36	99.96
37	101.29
38	100.82
39	99.54

(40, 7)

```
In [12]: # Removing < in "Percentage of Undernourishment Prevalence (3-Year Average)" column
# Because it changes in-place, we add try statements to prevent error
try:
    df["Percentage of Undernourishment Prevalence (3-Year Average)"] = pd.to_numeric(
except:
    pass
```

```
In [13]: # Making sure every dataset used are numbers
df = df.astype({'Percentage of Undernourishment Prevalence (3-Year Average)': 'float',
                'GDP per Capita (USD)': 'float',
                'Annual Inflation Rate (%)': 'float',
                'Food Production Index (2014-2016 = 100)': 'float'
                })
```

```
In [14]: display(df)
print(df.shape)
```

	Country	Year	\
0	Cambodia	2001	
1	Cambodia	2002	
2	Cambodia	2003	
3	Cambodia	2004	
4	Cambodia	2005	
5	Cambodia	2006	
6	Cambodia	2007	
7	Cambodia	2008	
8	Cambodia	2009	
9	Cambodia	2010	
10	Cambodia	2011	
11	Cambodia	2012	
12	Cambodia	2013	
13	Cambodia	2014	
14	Cambodia	2015	
15	Cambodia	2016	
16	Cambodia	2017	
17	Cambodia	2018	
18	Cambodia	2019	
19	Cambodia	2020	
20	Myanmar	2001	
21	Myanmar	2002	
22	Myanmar	2003	
23	Myanmar	2004	
24	Myanmar	2005	
25	Myanmar	2006	

26	Myanmar	2007
27	Myanmar	2008
28	Myanmar	2009
29	Myanmar	2010
30	Myanmar	2011
31	Myanmar	2012
32	Myanmar	2013
33	Myanmar	2014
34	Myanmar	2015
35	Myanmar	2016
36	Myanmar	2017
37	Myanmar	2018
38	Myanmar	2019
39	Myanmar	2020

	Percentage of Undernourishment Prevalence (3-Year Average) \
0	23.6
1	21.2
2	19.4
3	18.5
4	17.0
5	15.6
6	14.8
7	14.5
8	13.0
9	11.2
10	9.7
11	9.5
12	9.4
13	9.2
14	8.9
15	8.5
16	7.7
17	6.6
18	6.0
19	6.3
20	37.6
21	34.8
22	32.4
23	30.2
24	27.8
25	24.5
26	20.5
27	17.1
28	12.9
29	10.2
30	7.8
31	7.1

32	6.1
33	5.1
34	4.2
35	3.5
36	3.0
37	2.6
38	2.5
39	3.1

	Binary	Categorical	GDP per Capita (USD)	Annual Inflation Rate (%)	\
0		0	321.150224	-0.600648	
1		0	338.987477	0.211467	
2		0	362.335482	0.941746	
3		0	408.513639	4.319337	
4		0	474.111192	6.615259	
5		0	539.750329	5.810686	
6		0	631.525258	8.708828	
7		0	745.609127	24.096852	
8		0	738.054731	-1.241718	
9		0	785.502667	3.996395	
10		0	882.275614	5.478447	
11		0	950.880346	2.934316	
12		0	1013.420536	2.941625	
13		0	1093.495976	3.855689	
14		0	1162.904995	1.223932	
15		0	1269.591499	3.019140	
16		0	1385.260066	2.912636	
17		0	1512.126989	2.459085	
18		0	1643.121389	1.942575	
19		0	1547.511388	2.940295	
20		1	131.715298	21.101305	
21		1	128.099702	57.074511	
22		1	161.055524	36.589718	
23		1	193.368766	4.534214	
24		1	216.311501	9.368618	
25		1	240.624014	19.996487	
26		1	314.202294	35.024597	
27		1	460.908889	26.799537	
28		1	586.168180	1.472343	
29		1	746.945360	7.718382	
30		1	1061.344429	5.021460	
31		1	1134.302224	1.467583	
32		1	1168.165453	5.643039	
33		1	1210.097654	4.953299	
34		1	1196.743333	9.454172	
35		1	1136.610627	6.928825	
36		1	1151.114464	4.572537	
37		1	1250.173685	6.872329	

38	1	1271.111536	8.825067
39	1	1450.662673	7.092387

Food Production Index (2014-2016 = 100)

0	52.79
1	49.46
2	57.65
3	53.83
4	68.71
5	73.50
6	75.90
7	81.68
8	84.33
9	89.31
10	100.67
11	102.56
12	102.11
13	99.42
14	98.42
15	102.17
16	105.05
17	106.81
18	106.21
19	103.61
20	59.21
21	60.64
22	65.22
23	69.95
24	77.51
25	86.08
26	90.42
27	96.56
28	100.13
29	102.32
30	97.60
31	95.65
32	99.23
33	98.77
34	101.00
35	100.22
36	99.96
37	101.29
38	100.82
39	99.54

(40, 7)

```
In [15]: # List of all columns
print(list(df.columns))

['Country', 'Year', 'Percentage of Undernourishment Prevalence (3-Year Average)', 'Binary Cate
```

2.7 Functions

2.7.1 Preparation Functions

```
In [16]: # Preparation Functions
def get_features_targets(df, feature_names, target_names):
    df_feature = df.loc[:, feature_names]
    df_target = df.loc[:, target_names]
    return pd.DataFrame(df_feature), pd.DataFrame(df_target)

def split_data(df_feature, df_target, random_state=None, test_size=0.5):
    df_feature_rows, df_feature_columns = df_feature.shape
    array_all = list(range(0, df_feature_rows))

    np.random.seed(random_state)
    array_test = list(np.random.choice(array_all, int((df_feature_rows)*test_size), r
    array_train = [i for i in array_all if i not in array_test]

    df_feature_test = df_feature.iloc[array_test, :]
    df_feature_train = df_feature.iloc[array_train, :]
    df_target_test = df_target.iloc[array_test, :]
    df_target_train = df_target.iloc[array_train, :]

    return df_feature_train, df_feature_test, df_target_train, df_target_test

def normalize_z(dfin):
    mean = dfin.mean(axis=0)
    sd = dfin.std(axis=0)
    dfout = ((dfin.copy())-mean)/sd
    return dfout

def prepare_feature(df_feature):
    matrix_feature = (df_feature.copy()).to_numpy()
    matrix_one = np.ones([len(df_feature), 1])
    matrix_feature = np.concatenate((matrix_one, matrix_feature), axis=1)
    return matrix_feature

def prepare_target(df_target):
```

```
matrix_target = (df_target.copy()).to_numpy()
return matrix_target
```

2.7.2 Calculation Functions

In [17]: *# Calculation Functions*

```
def calc_linear(X, beta):
    beta_new_rows = int((X.size)/(len(X)))
    beta_new_columns = int((beta.size)/(beta_new_rows))
    beta = beta.reshape(beta_new_rows, beta_new_columns)
    return np.matmul(X, beta)

def compute_cost(X, y, beta):
    yhat = calc_linear(X, beta)
    yhat_y = yhat - y
    J = (np.matmul((yhat_y).T, yhat_y))/(2*len(X))
    return J

def gradient_descent(X, y, beta, alpha, num_iters):
    J_storage = np.array([])
    for i in range(num_iters):
        cost_value = calc_linear(X.T, ((calc_linear(X,beta))-y))/len(X)
        beta = beta - alpha*cost_value
        J_storage = np.append(J_storage, cost_value)
    return beta, J_storage

def predict(df_feature, beta):
    df_feature_z = normalize_z(df_feature.copy())
    X = prepare_feature(df_feature_z)
    yhat = calc_linear(X, beta)
    return yhat

def linear_regression(X, y, alpha, iterations):
    beta = np.zeros((X.shape[1], 1))
    beta, J_storage = gradient_descent(X, y, beta, alpha, iterations)
    yhat = predict(X, beta)
    return beta, J_storage, yhat
```

2.7.3 Metrics Functions

In [18]: *# Metrics Functions*

```
def r2_score(y, ypred):
    y_mean = np.mean(y)
    ss_tot = np.sum(np.power((y-y_mean), 2))
```

```

ss_res = np.sum(np.power(np.subtract(y, ypred), 2))
return 1 - ((ss_res)/ss_tot)

def mean_squared_error(y, ypred):
    ss_res = np.sum(np.power(np.subtract(y, ypred), 2))
    mse = (ss_res)/(len(y))
    return mse

def adjusted_r2_score(y, ypred, p):
    r2 = r2_score(y, ypred)
    n = y.shape[0]
    return 1 - (((1-r2)*(n-1))/(n-p-1))

def std_dev_score(y, ypred):
    n = y.shape[0]
    return (np.sum(np.subtract(y, ypred))/(n-1))**(0.5)

def std_error_reg_score(y, ypred, p):
    adjusted_r2 = adjusted_r2_score(y, ypred, p)
    std_dev = std_dev_score(y, ypred)
    return ((1-adjusted_r2)**(0.5))*std_dev

```

2.8 Features and Target Dataset Preparation

Describe here what are the features you use and why these features. Put any Python codes to prepare and clean up your features.

Do the same thing for the target. Describe your target and put any codes to prepare your target.

2.8.1 Choice of Features and Target

Features: By understanding the relation between the economic features: 1. GDP per Capita (USD) 2. Annual Inflation Rate (%) 3. Gross per capita Production Index Number (2014-2016 = 100) 4. Binary categorical

we can model them to predict prevalence of undernourishment. We use Binary Categorical as a feature where each country is represented by 0 or 1 (Cambodia:0, Myanmar:1) as a measure to prevent having to make two separate models for the two countries.

Target: We chose prevalence of undernourishment as a (%) as our target, as it gives us information about what percentage of people's food requirements are satisfied, encapsulating the aspect of food security.

```

In [19]: # DESCRIPTION:
        #'Country'

```

```

# 'Year'

# TARGET:
# 'Percentage of Undernourishment Prevalence (3-Year Average)'

# FEATURES:
# 'Binary Categorical'
# 'GDP per Capita (USD)'
# 'Annual Inflation Rate (%)'
# 'Gross per capita Production Index Number (2014-2016 = 100)'

```

2.9 Preparing Training and Test Dataset

```

In [20]: # Extract the features and the target
features = ['Binary Categorical', 'GDP per Capita (USD)', 'Annual Inflation Rate (%)']
targets = ['Percentage of Undernourishment Prevalence (3-Year Average)']
df_features, df_target = get_features_targets(df, features, targets) #DataFrame

# Split the data set into training and test
df_features_train, df_features_test, df_target_train, df_target_test = split_data(df_

# Normalize the features train using z normalization
df_features_train_z = normalize_z(df_features_train) #DataFrame

# Prepare the features train and target train to a NumPy
X = prepare_feature(df_features_train_z) #NumPy 5D
target = prepare_target(df_target_train) #NumPy 1D

In [21]: # Display the Descriptive Statistics
display(df_features.describe())
display(df_target.describe())

```

	Binary Categorical	GDP per Capita (USD)	Annual Inflation Rate (%)	\
count	40.00000	40.000000	40.000000	
mean	0.50000	825.396363	9.076909	
std	0.50637	454.494599	11.928724	
min	0.00000	128.099702	-1.241718	
25%	0.00000	396.969100	2.928896	
50%	0.50000	833.889141	4.987380	
75%	1.00000	1175.309923	8.737888	
max	1.00000	1643.121389	57.074511	

	Food Production Index (2014-2016 = 100)
count	40.000000
mean	87.907750
std	17.631598
min	49.460000
25%	75.300000

50%	98.010000
75%	100.865000
max	106.810000

	Percentage of Undernourishment Prevalence (3-Year Average)
count	40.000000
mean	13.590000
std	9.418607
min	2.500000
25%	6.525000
50%	9.950000
75%	18.725000
max	37.600000


2.9.1 Plotting Each Feature with Real Target (All Dataset)

Binary Categorical vs Percentage of Undernourishment Prevalence (3-Year Average)

In [22]: *# Scatter Plot of Binary Categorical vs Percentage of Undernourishment Prevalence (3-Year Average)*
ALL DATA SET

```
sns.set(rc={'figure.figsize':(9,6)})
myplot = sns.scatterplot(
    x = "Binary Categorical", y="Percentage of Undernourishment Prevalence",
    hue="Country", s=75, alpha=0.7)
myplot.legend(title = "Scatter Plot of " + "Binary Categorical" + " vs " + "Percentage of Undernourishment Prevalence",
    bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
    ncol=2, borderaxespad=0.)
```


Out[22]: <matplotlib.legend.Legend at 0x7f2897f27910>



GDP per Capita (USD) vs Percentage of Undernourishment Prevalence (3-Year Average)

```
In [23]: # Scatter Plot of GDP per Capita (USD) vs Percentage of Undernourishment Prevalence (3-Year Average)
# ALL DATA SET
sns.set(rc={'figure.figsize':(9,6)})
myplot = sns.scatterplot(
    x = "GDP per Capita (USD)", y="Percentage of Undernourishment Prevalence (3-Year Average)",
    hue="Country", s=75, alpha=0.7)
myplot.legend(title = "Scatter Plot of " + "GDP per Capita (USD)" + " vs " + "Percentage of Undernourishment Prevalence (3-Year Average)",
    bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
    ncol=2, borderaxespad=0.)
```

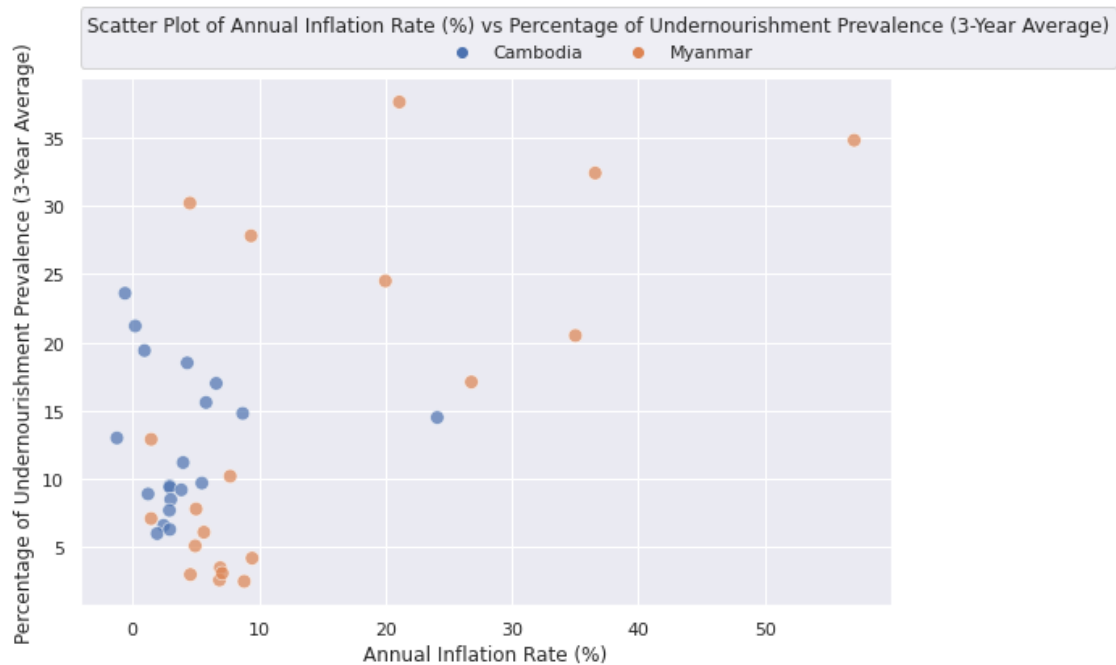
Out[23]: <matplotlib.legend.Legend at 0x7f289766ec90>



Annual Inflation Rate (%) vs Percentage of Undernourishment Prevalence (3-Year Average)

```
In [24]: # Annual Inflation Rate (%) vs Percentage of Undernourishment Prevalence (3-Year Average)
# ALL DATA SET
sns.set(rc={'figure.figsize':(9,6)})
myplot = sns.scatterplot(
    x = "Annual Inflation Rate (%)", y="Percentage of Undernourishment Prevalence (3-Year Average)",
    hue="Country", s=75, alpha=0.7)
myplot.legend(title = "Scatter Plot of " + "Annual Inflation Rate (%)" + " vs " + "Percentage of Undernourishment Prevalence (3-Year Average)",
    bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
    ncol=2, borderaxespad=0.)
```


Out[24]: <matplotlib.legend.Legend at 0x7f2895520290>



Food Production Index (2014-2016 = 100) vs Percentage of Undernourishment Prevalence (3-Year Average)

```
In [25]: # Food Production Index (2014-2016 = 100) vs Percentage of Undernourishment Prevalence
# ALL DATA SET
sns.set(rc={'figure.figsize':(9,6)})
myplot = sns.scatterplot(
    x = "Food Production Index (2014-2016 = 100)", y="Percentage of Undernourishment Prevalence (3-Year Average)",
    hue="Country", s=75, alpha=0.7)
myplot.legend(title = "Scatter Plot of " + "Food Production Index (2014-2016 = 100)" + " vs Percentage of Undernourishment Prevalence (3-Year Average)",
    bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
    ncol=2, borderaxespad=0.)
```

```
Out[25]: <matplotlib.legend.Legend at 0x7f289543e290>
```



2.10 Building Model

Describe your model. Is this Linear Regression or Logistic Regression? Put any other details about the model. Put the codes to build your model.

2.10.1 Calculating Initial Cost

$$J(\hat{\beta}_0, \hat{\beta}_1) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}(x^i) - y^i) \times (\hat{y}(x^i) - y^i)$$

```
In [26]: # Multiple Variables Cost Function
# Set the value of Beta (same size as features added by column vector of 1)
beta_multiple = np.zeros(((X.shape[1]), 1)) #NumPy 1D
J = compute_cost(X, target, beta_multiple)
print(J)
```

```
[[122.54625]]
```

2.10.2 Model Coefficients and Cost After Multiple Iterations

$$\begin{aligned}\hat{\beta}_0 &= \hat{\beta}_0 - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}(x^i) - y^i) x_0^i \\ \hat{\beta}_1 &= \hat{\beta}_1 - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}(x^i) - y^i) x_1^i \\ \hat{\beta}_2 &= \hat{\beta}_2 - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}(x^i) - y^i) x_2^i \\ &\dots \\ \hat{\beta}_n &= \hat{\beta}_n - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}(x^i) - y^i) x_n^i\end{aligned}$$

```
In [27]: # Beta After Iterations and J After Iterations
# Set the value of Iterations, Alpha, and Beta
```

```

alpha = 0.01
iterations = 1500
beta_multiple = np.zeros((X.shape[1]), 1))

# Call the gradient_descent function
beta_multiple, J_storage_multiple = gradient_descent(X, target, beta_multiple, alpha,
print(beta_multiple)

```

```

[[13.17499626]
 [ 0.35070389]
 [-5.13143412]
 [ 1.16119909]
 [-2.46487031]]

```

In [28]: *# Plot the graph of Cost Value in each iteration*

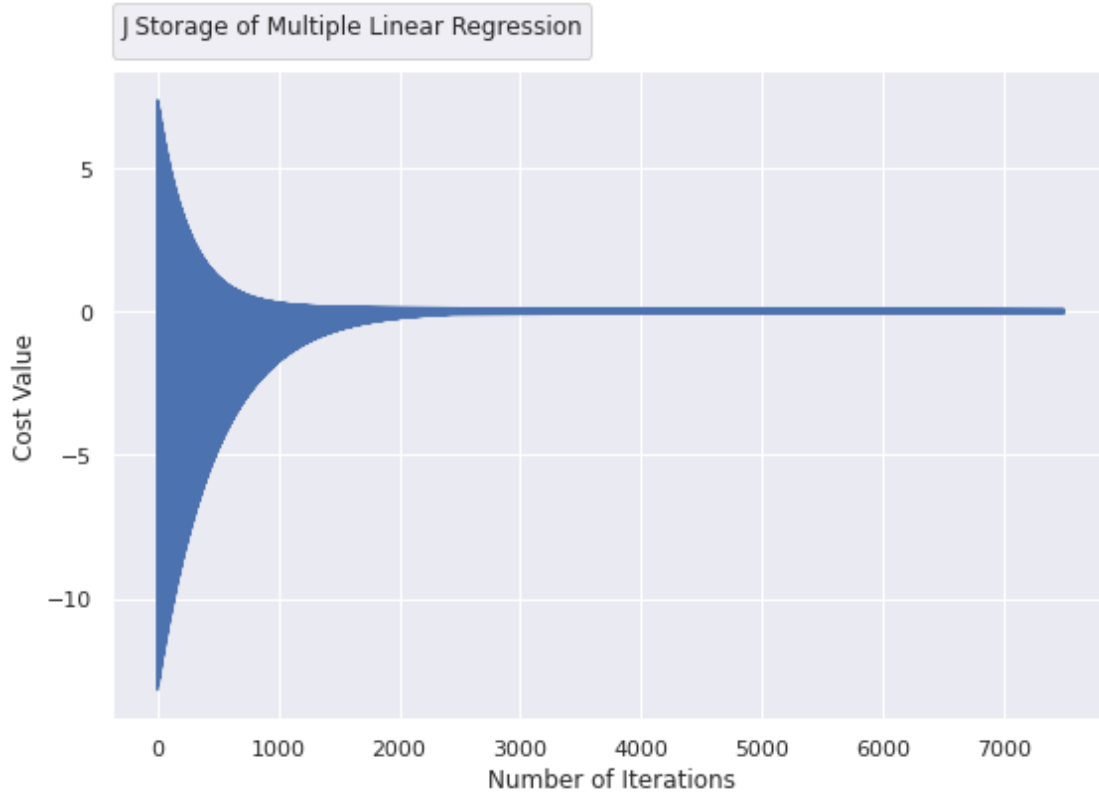
```

sns.set()
myplot = sns.lineplot(x=(np.linspace(start=0, stop=len(J_storage_multiple), num=len(J_
myplot.set_xlabel('Number of Iterations')
myplot.set_ylabel('Cost Value')
myplot.legend(title = "J Storage of Multiple Linear Regression",
               bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
               ncol=2, borderaxespad=0.)

```

No handles with labels found to put in legend.

Out[28]: <matplotlib.legend.Legend at 0x7f289548c790>



2.10.3 Intuitive Analysis of Model Coefficients

Intercept = 13.17

GDP per Capita (USD) coefficient: -5.13 (gdp increases, undernourishment decreases) indirectly proportional

Annual Inflation Rate (%) coefficient: 1.16 (inflation increases, undernourishment increases) directly proportional

Gross per capita Production Index Number (2014-2016 = 100) coefficient: -2.46 (agricultural production increases, undernourishment decreases) indirectly proportional

$$\mathbf{X} = \begin{bmatrix} 1 & x_1^1 & \dots & x_n^1 \\ 1 & x_1^2 & \dots & x_n^2 \\ \dots & \dots & \dots & \dots \\ 1 & x_1^m & \dots & x_n^m \end{bmatrix} \in \mathbb{R}^{m \times (n+1)}$$

$$\hat{\mathbf{b}} = \begin{bmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \\ \dots \\ \hat{\beta}_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$\hat{\mathbf{y}} = \mathbf{X} \times \hat{\mathbf{b}}$$

```
In [29]: # Predicted Value
          # Call the predict() method to get the predicted value of Feature Test
```

```

pred = predict(df_features_test, beta_multiple) #NumPy 1D
display(pred)

array([[20.79351752],
       [ 6.50075108],
       [10.99367766],
       [26.79523642],
       [12.82301609],
       [ 7.27461179],
       [20.21188228],
       [ 6.3503862 ],
       [29.50615341],
       [ 5.53880542],
       [16.76939911],
       [ 5.2657149 ],
       [ 5.25001583],
       [ 6.12764085],
       [21.90535443],
       [ 8.6937772 ]])

```

2.10.4 Improvement Iterations in Our Model

Prevalence of undernourishment has data from 2000 - 2021. Thus, we had 40 records for our two target countries. Due to our limited dataset, the diversity of training data played a huge role. We made efficient use of the limited data at hand to get the best possible adjusted r2 value and lowest standard error of regression by creating a function to find the best random seed value and therefore, the most diverse training data set. This in turn gave us a good adjusted r2 value and reduced standard error.

We also experimented with different ratios of training and test data sets.

```

In [30]: # ### finding best seed value with test_size=0.4 (After testing, we found that test_s
# ls_result = []
# max_result = 0
# seed_val = 0
# for val in range(999):
#     # Split the data set into training and test
#     df_features_train, df_features_test, df_target_train, df_target_test = split_da

#     # Normalize the features train using z normalization
#     df_features_train_z = normalize_z(df_features_train) #DataFrame

#     # Prepare the features train and target train to a NumPy
#     X = prepare_feature(df_features_train_z) #NumPy 5D
#     target = prepare_target(df_target_train) #NumPy 1D
#     # Beta After Iterations and J After Iterations
#     # Set the value of Iterations, Alpha
#     alpha = 0.01

```

```

#     iterations = 1500

#     # Call the gradient_descent function
#     beta_multiple, J_storage_multiple = gradient_descent(X, target, beta_multiple,
#     adjusted_r2 = adjusted_r2_score(y=prepare_target(df_target_test), ypred=pred, p
#     ls_result.append(adjusted_r2)
#     if max(ls_result)>max_result:
#         max_result = max(ls_result)
#         seed_val = val

# print("max adjusted r-squared:",max_result)
# print("corresponding seed value:", seed_val)

# RESULT
# max adjusted r-squared: 0.8742239115394154
# corresponding seed value: 99

```

2.10.5 Predicting Value of Target Using Trained Model on Test Dataset

```

In [31]: # Index used in test dataset after splitting data
df_plot_index = list(df_features_test.index)

```

```

In [32]: # Comparing Real Value of Test Dataset vs Predicted Value of Test Dataset
df_compare_target_predict = df.loc[df_plot_index, ["Country", "Year", "Percentage of U
df_compare_target_predict["Predicted Percentage of Undernourishment Prevalence (3-Year
df_compare_target_predict = df_compare_target_predict.sort_index(ascending=True)
display(df_compare_target_predict)
df_compare_target_predict = df_compare_target_predict.reset_index(drop=True)
df_compare_target_predict["Features"] = (np.linspace(1, len(df_compare_target_predict),

max_val_target_predict = df_compare_target_predict.loc[:, ["Percentage of Undernourishl
max_val_target_predict = max(max_val_target_predict.tolist())

```

	Country	Year	\
10	Cambodia	2011	
13	Cambodia	2014	
14	Cambodia	2015	
21	Myanmar	2002	
22	Myanmar	2003	
24	Myanmar	2005	
25	Myanmar	2006	
26	Myanmar	2007	
27	Myanmar	2008	
28	Myanmar	2009	
29	Myanmar	2010	
31	Myanmar	2012	
34	Myanmar	2015	
36	Myanmar	2017	

```

37 Myanmar 2018
38 Myanmar 2019

```

```

Percentage of Undernourishment Prevalence (3-Year Average) \
10 9.7
13 9.2
14 8.9
21 34.8
22 32.4
24 27.8
25 24.5
26 20.5
27 17.1
28 12.9
29 10.2
31 7.1
34 4.2
36 3.0
37 2.6
38 2.5

```

```

Predicted Percentage of Undernourishment Prevalence (3-Year Average)
10 8.693777
13 6.350386
14 5.538805
21 29.506153
22 26.795236
24 21.905354
25 20.793518
26 20.211882
27 16.769399
28 12.823016
29 10.993678
31 7.274612
34 6.127641
36 6.500751
37 5.265715
38 5.250016

```


```

In [33]: sns.set(rc={'figure.figsize':(9,6)})
sns.pointplot(data=df_compare_target_predict,
              x="Features", y="Predicted Percentage of Undernourishment Prevalence (3-Year Average)",
              color="orange", label="Real Value")
#plt.annotate("Real Value", (9.4, 8))
sns.pointplot(data=df_compare_target_predict,
              x="Features", y="Percentage of Undernourishment Prevalence (3-Year Average)",
              color="blue")

```

```
plt.annotate("Orange: Predicted Value",
            xy=(len(df_compare_target_predict)/1.311475, 1.0027*max_val_target_predict),
            color="orange")
plt.annotate("Blue: Real Value",
            xy=(len(df_compare_target_predict)/1.311475, max_val_target_predict/1.06),
            color="blue")
```

Out [33]: Text(12.200003812501192, 32.83018867924528, 'Blue: Real Value')



2.10.6 Combining Dataset for Easier Visualization

In [34]: # Make a new DataFrame for easier visualization

```
# 1st DataFrame for Real Value
df_plot_1 = df.loc[df_plot_index, :]
real_value_target = df_plot_1.pop("Percentage of Undernourishment Prevalence (3-Year Average)")
df_plot_1["Percentage of Undernourishment Prevalence (3-Year Average)"] = real_value_target
df_plot_1["Value Type"] = ""
df_plot_1.loc[:, "Value Type"] = "Real Value"

# 2nd DataFrame for Predicted Value
df_plot_2 = df.loc[df_plot_index, :]
df_plot_2.pop("Percentage of Undernourishment Prevalence (3-Year Average)")
predicted_value_target = pred
```



```

df_plot_2["Percentage of Undernourishment Prevalence (3-Year Average)"] = predicted_v
df_plot_2["Value Type"] = ""
df_plot_2.loc[:, "Value Type"] = "Predicted Value"

# Combine Together 1st DataFrame w/ 2nd DataFrame
df_plot = pd.DataFrame(pd.concat([df_plot_1, df_plot_2]))
df_plot = df_plot.sort_index(ascending=True)
display(df_plot)
print(df_plot.shape)

```

	Country	Year	Binary Categorical	GDP per Capita (USD) \
10	Cambodia	2011	0	882.275614
10	Cambodia	2011	0	882.275614
13	Cambodia	2014	0	1093.495976
13	Cambodia	2014	0	1093.495976
14	Cambodia	2015	0	1162.904995
14	Cambodia	2015	0	1162.904995
21	Myanmar	2002	1	128.099702
21	Myanmar	2002	1	128.099702
22	Myanmar	2003	1	161.055524
22	Myanmar	2003	1	161.055524
24	Myanmar	2005	1	216.311501
24	Myanmar	2005	1	216.311501
25	Myanmar	2006	1	240.624014
25	Myanmar	2006	1	240.624014
26	Myanmar	2007	1	314.202294
26	Myanmar	2007	1	314.202294
27	Myanmar	2008	1	460.908889
27	Myanmar	2008	1	460.908889
28	Myanmar	2009	1	586.168180
28	Myanmar	2009	1	586.168180
29	Myanmar	2010	1	746.945360
29	Myanmar	2010	1	746.945360
31	Myanmar	2012	1	1134.302224
31	Myanmar	2012	1	1134.302224
34	Myanmar	2015	1	1196.743333
34	Myanmar	2015	1	1196.743333
36	Myanmar	2017	1	1151.114464
36	Myanmar	2017	1	1151.114464
37	Myanmar	2018	1	1250.173685
37	Myanmar	2018	1	1250.173685
38	Myanmar	2019	1	1271.111536
38	Myanmar	2019	1	1271.111536

	Annual Inflation Rate (%)	Food Production Index (2014-2016 = 100) \
10	5.478447	100.67
10	5.478447	100.67
13	3.855689	99.42

13	3.855689	99.42
14	1.223932	98.42
14	1.223932	98.42
21	57.074511	60.64
21	57.074511	60.64
22	36.589718	65.22
22	36.589718	65.22
24	9.368618	77.51
24	9.368618	77.51
25	19.996487	86.08
25	19.996487	86.08
26	35.024597	90.42
26	35.024597	90.42
27	26.799537	96.56
27	26.799537	96.56
28	1.472343	100.13
28	1.472343	100.13
29	7.718382	102.32
29	7.718382	102.32
31	1.467583	95.65
31	1.467583	95.65
34	9.454172	101.00
34	9.454172	101.00
36	4.572537	99.96
36	4.572537	99.96
37	6.872329	101.29
37	6.872329	101.29
38	8.825067	100.82
38	8.825067	100.82

Percentage of Undernourishment Prevalence (3-Year Average) \

10	9.700000
10	8.693777
13	6.350386
13	9.200000
14	5.538805
14	8.900000
21	29.506153
21	34.800000
22	32.400000
22	26.795236
24	21.905354
24	27.800000
25	20.793518
25	24.500000
26	20.211882
26	20.500000
27	16.769399

27	17.100000
28	12.900000
28	12.823016
29	10.200000
29	10.993678
31	7.274612
31	7.100000
34	4.200000
34	6.127641
36	6.500751
36	3.000000
37	2.600000
37	5.265715
38	5.250016
38	2.500000

	Value Type
10	Real Value
10	Predicted Value
13	Predicted Value
13	Real Value
14	Predicted Value
14	Real Value
21	Predicted Value
21	Real Value
22	Real Value
22	Predicted Value
24	Predicted Value
24	Real Value
25	Predicted Value
25	Real Value
26	Predicted Value
26	Real Value
27	Predicted Value
27	Real Value
28	Real Value
28	Predicted Value
29	Real Value
29	Predicted Value
31	Predicted Value
31	Real Value
34	Real Value
34	Predicted Value
36	Predicted Value
36	Real Value
37	Real Value
37	Predicted Value
38	Predicted Value

38 Real Value


(32, 8)

2.10.7 Plotting Each Feature with Real Target and Predicted Target (All Dataset)

Binary Categorical vs Percentage of Undernourishment Prevalence (3-Year Average)

```
In [35]: # Scatter Plot of Binary Categorical vs Percentage of Undernourishment Prevalence (3-  
# TARGET DATASET  
sns.set(rc={'figure.figsize':(9,6)})  
myplot = sns.scatterplot(  
    x = "Binary Categorical", y="Percentage of Undernourishment Prevalence (3-Year Average)",  
    style="Value Type", hue="Country", s=75, alpha=0.7)  
myplot.legend(title = "Scatter Plot of " + "Binary Categorical" + " vs " + "Percentage of Undernourishment Prevalence (3-Year Average)",  
    bbox_to_anchor=(0., 1.02, 1., .102), loc=3,  
    ncol=2, borderaxespad=0.)
```


Out [35]: <matplotlib.legend.Legend at 0x7f289528ac10>



GDP per Capita (USD) vs Percentage of Undernourishment Prevalence (3-Year Average)

```
In [36]: # Scatter Plot of GDP per Capita (USD) vs Percentage of Undernourishment Prevalence (3-Year Average)
# TARGET DATASET
sns.set(rc={'figure.figsize':(9,6)})
myplot = sns.scatterplot(
    x = "GDP per Capita (USD)", y="Percentage of Undernourishment Prevalence (3-Year Average)",
    style="Value Type", hue="Country", s=75, alpha=0.7)
myplot.legend(title = "Scatter Plot of " + "GDP per Capita (USD)" + " vs " + "Percentage of Undernourishment Prevalence (3-Year Average)",
    bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
    ncol=2, borderaxespad=0.)
```


Out[36]: <matplotlib.legend.Legend at 0x7f2895035350>



Annual Inflation Rate (%) vs Percentage of Undernourishment Prevalence (3-Year Average)

```
In [37]: # Scatter Plot of Annual Inflation Rate (%) vs Percentage of Undernourishment Prevalence (3-Year Average)
# TARGET DATASET
sns.set(rc={'figure.figsize':(9,6)})
myplot = sns.scatterplot(
    x = "Annual Inflation Rate (%)", y="Percentage of Undernourishment Prevalence (3-Year Average)",
    style="Value Type", hue="Country", s=75, alpha=0.7)
myplot.legend(title = "Scatter Plot of " + "Annual Inflation Rate (%)" + " vs " + "Percentage of Undernourishment Prevalence (3-Year Average)",
    bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
    ncol=2, borderaxespad=0.)
```


Out [37]: <matplotlib.legend.Legend at 0x7f2895020b90>



Food Production Index (2014-2016 = 100) vs Percentage of Undernourishment Prevalence (3-Year Average)

```
In [38]: # Scatter Plot of Food Production Index (2014-2016 = 100) vs Percentage of Undernourishment Prevalence (3-Year Average)
# TARGET DATASET
sns.set(rc={'figure.figsize':(9,6)})
myplot = sns.scatterplot(
    x = "Food Production Index (2014-2016 = 100)", y="Percentage of Undernourishment Prevalence (3-Year Average)",
    style="Value Type", hue="Country", s=75, alpha=0.7)
myplot.legend(title = "Scatter Plot of " + "Food Production Index (2014-2016 = 100)" + " vs Percentage of Undernourishment Prevalence (3-Year Average)",
    bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
    ncol=2, borderaxespad=0.)
```

Out [38]: <matplotlib.legend.Legend at 0x7f2894eace10>



2.11 Evaluating the Model

Describe your metrics and how you want to evaluate your model. Put any Python code to evaluate your model. Use plots to have a visual evaluation.

2.11.1 Standard Metrics

Mean Squared Error (MSE) $MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$

```
In [39]: # # Mean Squared Error (MSE)
# mse = mean_squared_error(y=prepare_target(df_target_test), ypred=pred)
# print(mse)
# Mean Squared Error (MSE)
mse = mean_squared_error(prepare_target(df_target_test), pred)
print(mse)
```

9.990736660164742

Coefficient of Determination (R²) $R^2 = 1 - \frac{SS_{RES}}{SS_{TOT}} = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$

```
In [40]: # Coefficient of Determination (R2)
# Calculate r2 score by calling a function, the arguments must be a NumPy
r2 = r2_score(prepare_target(df_target_test), pred)
print(r2)
```

0.9077642017955713

2.11.2 Relevant Metrics

Adjusted Coefficient of Determination (Adjusted R²) We cannot use R² to evaluate multiple linear regression due to the following issues 1. Every time we add a predictor to a model, the R-squared increases, even if due to chance alone. It never decreases. Consequently, a model with more terms may appear to have a better fit simply because it has more terms. 2. If a model has too many predictors and higher order polynomials, it begins to model the random noise in the data. This leads to overfitting the model and it produces misleadingly high R-squared values and a lessened ability to make predictions.

Adjusted R-squared is a modified version of R-squared that has been adjusted for the number of predictors in the model. Adjusted R-squared helps us determine how much of the correlation with the index is due to the addition of those variables. The adjusted R-squared compensates for the addition of variables and only increases if the new predictor enhances the model. (Side note: The Adjusted R2 Score is an extension of the R2 Score that takes into account sample size and the number of independent variables. Hence, it is possible for it to return zero or negative values in the event of a lack of data when doing regression)

$$R_{adj}^2 = 1 - \frac{(1-R^2)(n-1)}{n-p-1} \quad n : \text{number of dataset } p : \text{number of features}$$

```
In [41]: # Adjusted Coefficient of Determination (Adjusted R2)
# Calculate adjusted r2 score by calling a function, the arguments must be a NumPy for
# p is the number of Independent Variables, which are the length of Features list
adjusted_r2 = adjusted_r2_score(y=prepare_target(df_target_test), ypred=pred, p=len(fe
print(adjusted_r2)
```

0.8742239115394154

Standard Error of Regression (S) In addition to the adjusted R-squared metric, we also used standard error or regression.

The standard error of the regression (S), represents the average distance that the observed values fall from the regression line. S is in the units of the dependent variable.

It is particularly useful because we can use it to assess the precision of predictions. Roughly 95% of the observation should fall within +/- two standard errors of the regression, which is a quick approximation of a 95% prediction interval.

$$S = \sqrt{1 - R_{adj}^2} \cdot S.D$$
$$S.D = \sqrt{\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n-1}}$$

```
In [42]: # Standard Error of Regression (S)
# p is the number of Independent Variables, which are the length of Features list
std_eror_reg = std_error_reg_score(y=prepare_target(df_target_test), ypred=pred, p=len(fe
print(std_eror_reg)
```

0.3730853869367281

2.12 Verifying Process with Scikit Sklearn

```
In [43]: from sklearn.linear_model import LinearRegression
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import r2_score, mean_squared_error
```

```
In [44]: # Extract the features and the target
        df_features_scikit, df_target_scikit = get_features_targets(df, ['Binary Categorical'])

        # Normalize the features train using z normalization
        df_features_scikit = normalize_z(df_features_scikit)

        # Split the data into training and test data set using scikit-learn function
        df_features_train_scikit, df_features_test_scikit, df_target_train_scikit, df_target_test_scikit = train_test_split(df_features_scikit, df_target_scikit, test_size=0.2, random_state=42)

        # Instantiate LinearRegression() object
        model = LinearRegression()

        # Call the fit() method and find the Beta Value
        model.fit(df_features_train_scikit, df_target_train_scikit)
        print(model.coef_, model.intercept_)

        # Call the predict() method
        pred_scikit = model.predict(df_features_test_scikit)

[[ 0.32477307 -5.66203628  2.25446523 -1.60300413]] [14.3862165]
```

```
In [45]: # Call the r2_score method and find the r2 value
        print(r2_score(df_target_test_scikit, pred_scikit))

0.8759086808324694
```

2.13 Improving the Model

Discuss any steps you can do to improve the models. Put any python codes. You can repeat the steps above with the codes to show the improvement in the accuracy.

- Adding more dataset to improve precision of the model
 1. More countries
 2. More years
- Adding more features that is relevant to improve accuracy of the model
- Adding more targets to better analyse and predict food security and food safety
- Adding interactive window with input() function to allow one to choose which features and targets one wants to model and select what types of graphs one wishes to see
- Using more Python library such as TensorFlow to get the derivatives of the Cost Function, instead of using Gradient Descent

2.14 Discussion and Analysis

Discuss your model and accuracy in solving the problem. Analyze the results of your metrics. Put any conclusion here.

Our model has pretty good accuracy in predicting Percentage of Prevalence of Undernourishment with adjusted r-squared value = 0.8742239115394154 and Standard error of regression (S) = 0.3730853869367281. Percentage of Prevalence of Undernourishment gives us information about the food security of the country. With this data we can predict food security of the country in the future and take necessary steps and prepare in advance to avoid low levels of food security.

In []: